



MNIST Dataset

The MNIST database of handwritten digits (<http://yann.lecun.com>)

[Data Card](#)[Code \(225\)](#)[Discussion \(0\)](#)[Suggestions \(0\)](#)

About Dataset

Context

MNIST is a subset of a larger set available from NIST (it's copied from <http://yann.lecun.com/exdb/mnist/>)

Content

The MNIST database of handwritten digits has a training set of 60,000 examples, and a test set of 10,000 examples. .

Four files are available:

- train-images-idx3-ubyte.gz: training set images (9912422 bytes)
- train-labels-idx1-ubyte.gz: training set labels (28881 bytes)
- t10k-images-idx3-ubyte.gz: test set images (1648877 bytes)
- t10k-labels-idx1-ubyte.gz: test set labels (4542 bytes)

How to read

See [sample MNIST reader](#)

Usability ⓘ

7.50

License


Data files © Original Authors

Expected update frequency

Not specified

Tags

[Computer Science](#)[Image](#)[Multiclass Classification](#)



Trianing Data

- 60,000 images of handwritten digits from 250 individuals. Among these, 50% are high school students and the other 50% are employees from the Census Bureau.
- All images are composed of 28×28 pixels. Each image also has a corresponding label indicating the digit represented in the image.

Testing Data


- The test data contains 10,000 images of handwritten digits, also sourced from U.S. high school students and Census Bureau employees.
- These images come from a different set of 250 individuals than those whose handwriting was included in the training data.

Why MNIST Dataset?

- **Data Availability:** The MNIST dataset, which stands for the Modified National Institute of Standards and Technology dataset, is widely used for training and testing in the field of machine learning. It is a large database of handwritten digits that is publicly available and thus easy to access.
- **Well-defined Problem:** Since the dataset and problem (digit recognition) are well-defined, you can focus more on the modeling aspects rather than spending a lot of time in the data gathering and cleaning stages.


Identify a Deep Learning Problem

- The MNIST dataset is ideal for applying CNN, which are a class of deep neural networks highly effective at recognizing patterns and structures in images like the digits in MNIST.



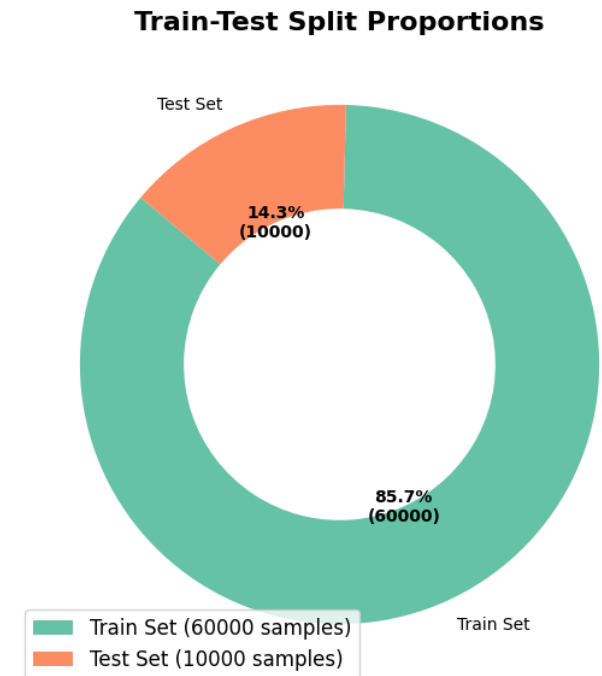
EDA

- We take a look into the input data
- Most of the handwritings are easy to identify by human eyes
- But I do see a variety of writing styles




EDA

- Each input picture is 224x224x3
- Dataset is well balanced among all 10 digits
- Pre-defined test/train split
- I believe no transform is required, as 2D pictures are ideal target of CNN




Analysis: Deep-Learning model

- Since 224x224 is the input size of the famous VGG16, we use the pre-trained model provided in keras to solve the problem.
- Below is the architecture of VGG16. The schematic consists of CNN feature extractor and the following fully-connected classifier.



Analysis: Deep-Learning model Performance

- Achieved 99.5% accuracy on validation set.
- Accuracy is almost 90% at the beginning of training, thanks to the pretrained weight.
- Run time is ~15min



Analysis: Unsupervised model

- VGG is a bit overkill for this problem
- Unsupervised methods are usually less powerful but faster.
 - Trade accuracy for shorter run time
- We use TSNE to perform dimension reduction
- Followed by GMM clustering
- Ended with label permutation to match with ground truth

```
X_train_tsne=X_train.reshape(4800,-1)
print(X_train_tsne.shape)
tsne = TSNE(n_components=3,
            n_jobs=-1)
digits_proj = tsne.fit_transform(X_train_tsne)
gmm = GaussianMixture(
    n_components=10)
clusters = gmm.fit_predict(digits_proj)
```

Analysis: Unsupervised model Performance

- The accuracy drops to 83.6%
- Run time reduce to 2 min
- Increasing TSNE component does not improve accuracy

```
stop = timeit.default_timer()

print('Time: ', stop - start)
```

```
(4800, 224, 224, 3)
(4800, 150528)
Time: 157.55008647299974
```

```
voter=np.zeros(shape=(10,10))
y_train_label=np.argmax(y_train,axis=1)
for i in range(len(clusters)):
    #print(clusters[i],y_train_label[i])
    voter[clusters[i]][y_train_label[i]]+=1
print(voter)
trans=np.argmax(voter,axis=1)
correct=0
for i in range(len(clusters)):
    if y_train_label[i]==trans[clusters[i]]:
        correct+=1
print(correct/len(clusters))
```

```
[[ 1.  1.  6. 448.  0. 29.  0.  0. 33.  7.]
 [ 0.  3.  9.  4. 13.  0.  0. 470.  1. 159.]
 [ 1.  0.  3.  2. 10.  7. 479.  0.  3.  1.]
 [476.  0.  1.  1.  0.  1.  3.  0.  7.  5.]
 [ 0. 206.  8.  3.  4.  1.  1.  4.  5.  0.]
 [ 0.  2. 429.  1.  1.  1.  0.  4.  2.  0.]
 [ 0.  1.  2.  7.  0.  1.  0.  0. 347.  0.]
 [ 5.  1.  3.  8.  0. 366.  4.  1. 32.  0.]
 [ 0.  1.  1.  3. 472. 10.  1. 25.  8. 299.]
 [ 1. 322.  4.  4.  4.  1.  1.  5.  2.  2.]]
0.8364583333333333
```

Conclusion

- Pre-trained model facilitates model convergence
- CNN is more powerful than machine learning methods regarding 2D image classification.
- CNN are more time-consuming than machine learning methods.

```

import numpy as np
import pandas as pd

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

/kaggle/input/mnist-dataset/train-images.idx3-ubyte
/kaggle/input/mnist-dataset/t10k-labels.idx1-ubyte
/kaggle/input/mnist-dataset/t10k-images.idx3-ubyte
/kaggle/input/mnist-dataset/train-labels.idx1-ubyte
/kaggle/input/mnist-dataset/t10k-labels-idx1-ubyte/t10k-labels-idx1-
ubyte
/kaggle/input/mnist-dataset/t10k-images-idx3-ubyte/t10k-images-idx3-
ubyte
/kaggle/input/mnist-dataset/train-labels-idx1-ubyte/train-labels-idx1-
ubyte
/kaggle/input/mnist-dataset/train-images-idx3-ubyte/train-images-idx3-
ubyte

```

Libraries to be used

```

import tensorflow as tf
from keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten,
Dense, Dropout
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

```

Load Data

```

(train_images, train_labels), (test_images, test_labels) =
mnist.load_data()
train_images, test_images = train_images / 255.0, test_images / 255.0
# normalising the dimensions of images to increase the performance of
cnn

# Load the MNIST dataset
(train_images, train_labels), (test_images, test_labels) =
mnist.load_data()

# Normalize the images
train_images, test_images = train_images / 255.0, test_images / 255.0

```

```

# Use only 60% of the dataset
train_images, train_labels = train_images[:int(0.1 *
len(train_images))], train_labels[:int(0.1 * len(train_labels))]

# Convert images to 3-channel (RGB) and resize to 224x224 for VGG16
def preprocess_images(images):
    images = tf.expand_dims(images, axis=-1) # Add channel dimension
    images = tf.image.grayscale_to_rgb(images) # Convert to RGB
    images = tf.image.resize(images, [224, 224]) # Resize to 224x224
    return images.numpy()

train_images = preprocess_images(train_images)
test_images = preprocess_images(test_images)

# Convert labels to one-hot encoding
train_labels = to_categorical(train_labels, num_classes=10)
test_labels = to_categorical(test_labels, num_classes=10)

# Split train data into train and validation sets (80-20 split of 60%
dataset)
X_train, X_val, y_train, y_val = train_test_split(train_images,
train_labels, test_size=0.2, random_state=42)

random_index = np.random.randint(0, len(train_images))
random_image = train_images[random_index]
random_label = train_labels[random_index]

```

EDA

```

temp=np.argmax(train_labels, axis=1)
d={}
for i in range(10):
    d[i]=0
for q in temp:
    d[q]+=1
print(d)
#balanced dataset

{0: 592, 1: 671, 2: 581, 3: 608, 4: 623, 5: 514, 6: 608, 7: 651, 8:
551, 9: 601}

for i in range(20):
    random_index = np.random.randint(0, len(train_images))
    random_image = train_images[random_index]
    random_label = train_labels[random_index]

    plt.imshow(random_image, cmap='gray') # Use 'gray' for MNIST
images

```

```
plt.axis('off')  
plt.title(np.argmax(random_label)) # Convert one-hot label back  
to a class index  
plt.show()
```

1



0



6



9



4



1



3



1



8



7



0



4



4



8



9



9



6



0



7



8



```
random_image.shape
(224, 224, 3)

labels = str(list(range(0,10)))
labels
'[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]'

num_samples = 12
random_indices = np.random.choice(train_images.shape[0], num_samples,
replace=False)
sample_images = train_images[random_indices]
sample_labels = train_labels[random_indices]

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Flatten, Dense, Dropout
vgg_base = VGG16(weights='imagenet', include_top=False,
input_shape=(224, 224, 3))

# Freeze the VGG16 base layers
for layer in vgg_base.layers:
    layer.trainable = False
```

```
# Add custom layers on top of VGG16
x = Flatten()(vgg_base.output)
x = Dense(128, activation='relu')(x)
x = Dropout(0.5)(x)
output = Dense(10, activation='softmax')(x) # 10 classes for MNIST
```

```
# Create final model
model = Model(inputs=vgg_base.input, outputs=output)
# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58889256/58889256 _____ 0s 0us/step
```

```
# Train the model without using ImageDataGenerator
history = model.fit(
    X_train, y_train, # Directly use the training data
    validation_data=(X_val, y_val), # Directly use the validation
data
    epochs=10,
    batch_size=32 # Set batch size here
)
```

Epoch 1/10

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

```
I0000 00:00:1733860691.609322      104 service.cc:145] XLA service
0x7bb580105750 initialized for platform CUDA (this does not guarantee
that XLA will be used). Devices:
I0000 00:00:1733860691.609369      104 service.cc:153]   StreamExecutor
device (0): Tesla T4, Compute Capability 7.5
I0000 00:00:1733860691.609373      104 service.cc:153]   StreamExecutor
device (1): Tesla T4, Compute Capability 7.5
```

```
1/150 _____ 43:32 18s/step - accuracy: 0.0938 - loss:
2.5070
```

```
I0000 00:00:1733860707.849357      104 device_compiler.h:188] Compiled
cluster using XLA! This line is logged at most once for the lifetime
of the process.
```

```
150/150 _____ 57s 262ms/step - accuracy: 0.3393 - loss:
1.9881 - val_accuracy: 0.9492 - val_loss: 0.4647
```

Epoch 2/10

```
150/150 _____ 26s 176ms/step - accuracy: 0.6187 - loss:
0.9654 - val_accuracy: 0.9700 - val_loss: 0.2271
```

Epoch 3/10

```

150/150 _____ 25s 166ms/step - accuracy: 0.6688 - loss:
0.8045 - val_accuracy: 0.9700 - val_loss: 0.1801
Epoch 4/10
150/150 _____ 26s 171ms/step - accuracy: 0.7079 - loss:
0.7316 - val_accuracy: 0.9725 - val_loss: 0.1468
Epoch 5/10
150/150 _____ 26s 173ms/step - accuracy: 0.7163 - loss:
0.6900 - val_accuracy: 0.9700 - val_loss: 0.1375
Epoch 6/10
150/150 _____ 25s 170ms/step - accuracy: 0.7196 - loss:
0.6797 - val_accuracy: 0.9750 - val_loss: 0.1097
Epoch 7/10
150/150 _____ 26s 170ms/step - accuracy: 0.7473 - loss:
0.6259 - val_accuracy: 0.9775 - val_loss: 0.0926
Epoch 8/10
150/150 _____ 26s 172ms/step - accuracy: 0.7272 - loss:
0.6564 - val_accuracy: 0.9742 - val_loss: 0.1076
Epoch 9/10
150/150 _____ 26s 172ms/step - accuracy: 0.7298 - loss:
0.6290 - val_accuracy: 0.9833 - val_loss: 0.0817
Epoch 10/10
150/150 _____ 26s 171ms/step - accuracy: 0.7240 - loss:
0.6339 - val_accuracy: 0.9683 - val_loss: 0.1314

```

Unsupervised method: TSNE+GMM

```

from sklearn.manifold import Isomap, TSNE
from sklearn.decomposition import PCA
# Clustering
from sklearn.mixture import GaussianMixture
import timeit

start = timeit.default_timer()

print(X_train.shape)
X_train_tsne=X_train.reshape(4800,-1)
print(X_train_tsne.shape)
tsne = TSNE(n_components=3,
            n_jobs=-1)
digits_proj = tsne.fit_transform(X_train_tsne)
gmm = GaussianMixture(
    n_components=10)
clusters = gmm.fit_predict(digits_proj)
#Your statements here

stop = timeit.default_timer()

print('Time: ', stop - start)

```

```

(4800, 224, 224, 3)
(4800, 150528)
Time: 157.55008647299974

voter=np.zeros(shape=(10,10))
y_train_label=np.argmax(y_train,axis=1)
for i in range(len(clusters)):
    #print(clusters[i],y_train_label[i])
    voter[clusters[i]][y_train_label[i]]+=1
print(voter)
trans=np.argmax(voter,axis=1)
correct=0
for i in range(len(clusters)):
    if y_train_label[i]==trans[clusters[i]]:
        correct+=1
print(correct/len(clusters))

[[ 1.  1.  6. 448.  0. 29.  0.  0. 33.  7.]
 [ 0.  3.  9.  4. 13.  0.  0. 470.  1. 159.]
 [ 1.  0.  3.  2. 10.  7. 479.  0.  3.  1.]
 [476.  0.  1.  1.  0.  1.  3.  0.  7.  5.]
 [ 0. 206.  8.  3.  4.  1.  1.  4.  5.  0.]
 [ 0.  2. 429.  1.  1.  1.  0.  4.  2.  0.]
 [ 0.  1.  2.  7.  0.  1.  0.  0. 347.  0.]
 [ 5.  1.  3.  8.  0. 366.  4.  1. 32.  0.]
 [ 0.  1.  1.  3. 472. 10.  1. 25.  8. 299.]
 [ 1. 322.  4.  4.  4.  1.  1.  5.  2.  2.]]
0.8364583333333333

```

Unsupervised method: reduce training data

```

start = timeit.default_timer()

X_train_tsne2=X_train_tsne[:2400][:]
y_train_label2=np.argmax(y_train,axis=1)[:2400]
print(X_train_tsne2.shape)
tsne2 = TSNE(n_components=3,
             n_jobs=-1)
digits_proj2 = tsne2.fit_transform(X_train_tsne2)
gmm2 = GaussianMixture(
    n_components=10)
clusters2 = gmm2.fit_predict(digits_proj2)
#Your statements here

stop = timeit.default_timer()

print('Time: ', stop - start)

```

```

voter2=np.zeros(shape=(10,10))

for i in range(len(clusters2)):
    #print(clusters[i],y_train_label[i])
    voter2[clusters2[i]][y_train_label2[i]]+=1
print(voter2)
trans2=np.argmax(voter2,axis=1)
correct2=0
for i in range(len(clusters2)):
    if y_train_label2[i]==trans2[clusters2[i]]:
        correct2+=1
print(correct2/len(clusters2))

(2400, 150528)
Time: 54.03585478500008
[[ 1.  2.  0.  4. 243.  7.  2. 10.  3. 174.]
 [ 0.  0.  5.  7.  0. 170.  4.  1.  2.  0.]
 [232.  0.  0.  0.  0.  0.  2.  0.  1.  1.]
 [ 0.  1.  5. 223.  0. 23.  0.  0. 10.  3.]
 [ 1. 147.  3.  0.  0.  1.  0.  5.  1.  0.]
 [ 0.  2.  9.  1. 13.  0.  0. 220.  0. 41.]
 [ 0.  1. 218.  6.  0.  0.  0.  3.  5.  0.]
 [ 0.  0.  1.  2.  6.  6. 241.  0.  4.  0.]
 [ 0. 110.  3.  2.  2.  0.  0.  3.  2.  0.]
 [ 0.  0.  0.  7.  0.  4.  0.  0. 194.  0.]]
0.8325

```

```

from tensorflow.keras.optimizers import Adagrad

# Modeling
model.compile(optimizer='Adagrad', loss='categorical_crossentropy',
metrics=['accuracy'])

model.summary()

Model: "functional_1"

```

Layer (type)	Output Shape	
Param #		
input_layer (InputLayer)	(None, 224, 224, 3)	
0		
block1_conv1 (Conv2D)	(None, 224, 224, 64)	
1,792		

36,928	block1_conv2 (Conv2D)	(None, 224, 224, 64)
0	block1_pool (MaxPooling2D)	(None, 112, 112, 64)
73,856	block2_conv1 (Conv2D)	(None, 112, 112, 128)
147,584	block2_conv2 (Conv2D)	(None, 112, 112, 128)
0	block2_pool (MaxPooling2D)	(None, 56, 56, 128)
295,168	block3_conv1 (Conv2D)	(None, 56, 56, 256)
590,080	block3_conv2 (Conv2D)	(None, 56, 56, 256)
590,080	block3_conv3 (Conv2D)	(None, 56, 56, 256)
0	block3_pool (MaxPooling2D)	(None, 28, 28, 256)
1,180,160	block4_conv1 (Conv2D)	(None, 28, 28, 512)
2,359,808	block4_conv2 (Conv2D)	(None, 28, 28, 512)
2,359,808	block4_conv3 (Conv2D)	(None, 28, 28, 512)

0	block4_pool (MaxPooling2D)	(None, 14, 14, 512)	
	block5_conv1 (Conv2D)	(None, 14, 14, 512)	
2,359,808			
	block5_conv2 (Conv2D)	(None, 14, 14, 512)	
2,359,808			
	block5_conv3 (Conv2D)	(None, 14, 14, 512)	
2,359,808			
0	block5_pool (MaxPooling2D)	(None, 7, 7, 512)	
0	flatten (Flatten)	(None, 25088)	
	dense (Dense)	(None, 128)	
3,211,392			
0	dropout (Dropout)	(None, 128)	
	dense_1 (Dense)	(None, 10)	
1,290			

Total params: 17,927,370 (68.39 MB)

Trainable params: 3,212,682 (12.26 MB)


Non-trainable params: 14,714,688 (56.13 MB)

```
loss,accuracy = model.evaluate(X_train,y_train)
print(f"The model accuracy is : {accuracy} \n the model loss : {loss}")
```


150/150 ————— 21s 138ms/step - accuracy: 0.9898 - loss: 0.0684

The model accuracy is : 0.9879166483879089
the model loss : 0.07202592492103577

```
plt.plot(history.history['accuracy'],label='Training Accuracy')  
plt.plot(history.history['val_accuracy'], label = 'Validation  
Accuracy')  
plt.xlabel("Epochs")  
plt.ylabel('Accuracy')  
plt.legend()  
plt.show()
```



```
plt.plot(history.history['loss'],label='Training Loss')  
plt.plot(history.history['val_loss'],label='Validation Loss')  
plt.xlabel("Epoche")  
plt.ylabel("Loss")  
plt.legend()  
plt.show()
```



```
import numpy as np

# Predict the class probabilities for the test dataset
predictions = model.predict(X_val)

# Convert probabilities to class labels
predicted_labels = np.argmax(predictions, axis=1)
true_labels = np.argmax(y_val, axis=1) # Convert one-hot to class indices

38/38 ————— 6s 150ms/step

from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Generate the confusion matrix
true_labels = np.argmax(y_val, axis=1) # Convert one-hot encoded to class indices
predicted_labels = np.argmax(predictions, axis=1) # Already converting probabilities to class indices
correct=0
for i in range(len(true_labels)):
    if true_labels[i]==predicted_labels[i]:
        correct+=1
```

```

print(correct/len(true_labels))
from sklearn.metrics import confusion_matrix

# Generate confusion matrix
conf_matrix = confusion_matrix(true_labels, predicted_labels)

0.9683333333333334

# Plot the confusion matrix
plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=range(10), yticklabels=range(10))
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()

```

